




# Convolutional Neural Networks




# Convolutional Neural Networks

## Overview

- › special architecture of deep neural networks
- › Computer Vision
- › automated feature extraction
- › classification of images through automatically extracted features

# Convolutional Neural Networks

## Overview



# Convolutional Neural Networks


## Overview

- › Convolutional Layer → Feature extraction
- › Pooling Layer → Dimensionality reduction
- › Fully Connected → Classification

# Convolutional Neural Networks

## Convolutional Layer

32x32x3 image



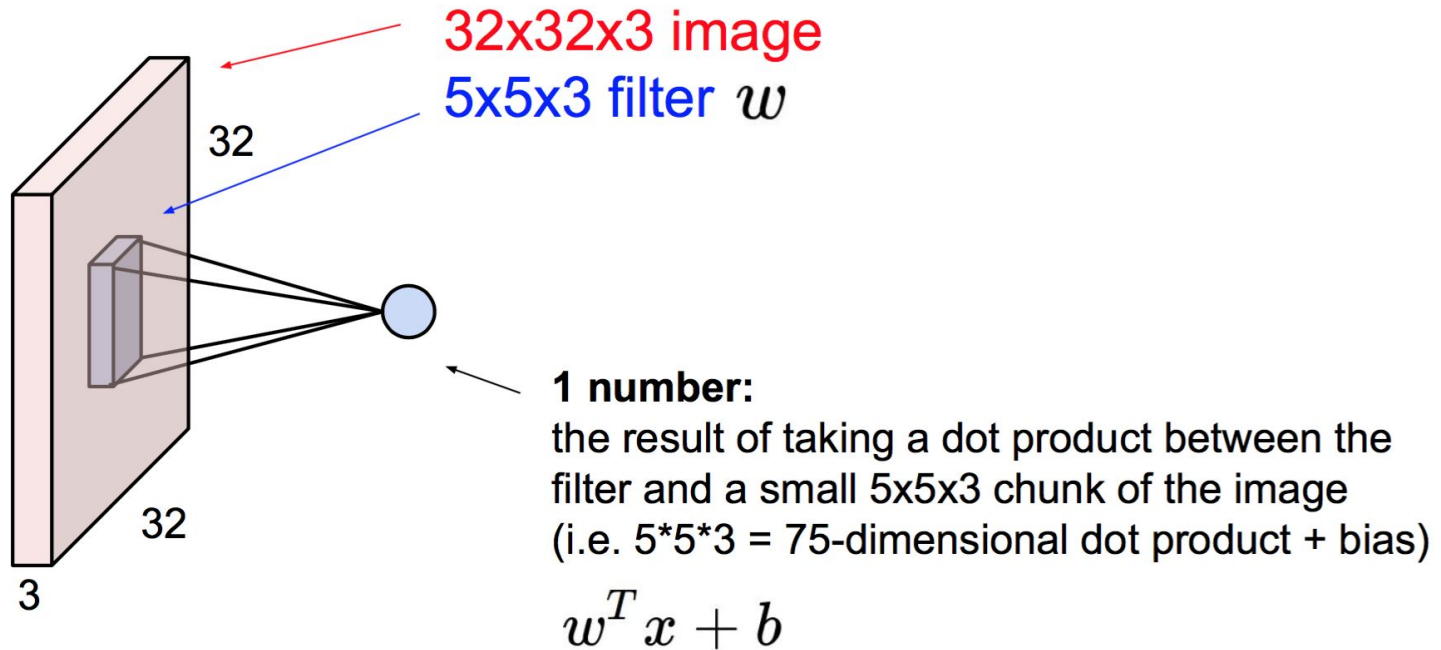
5x5x3 filter



**Convolve** the filter with the image  
i.e. “slide over the image spatially,  
computing dot products”


# Convolutional Neural Networks

## Convolutional Layer



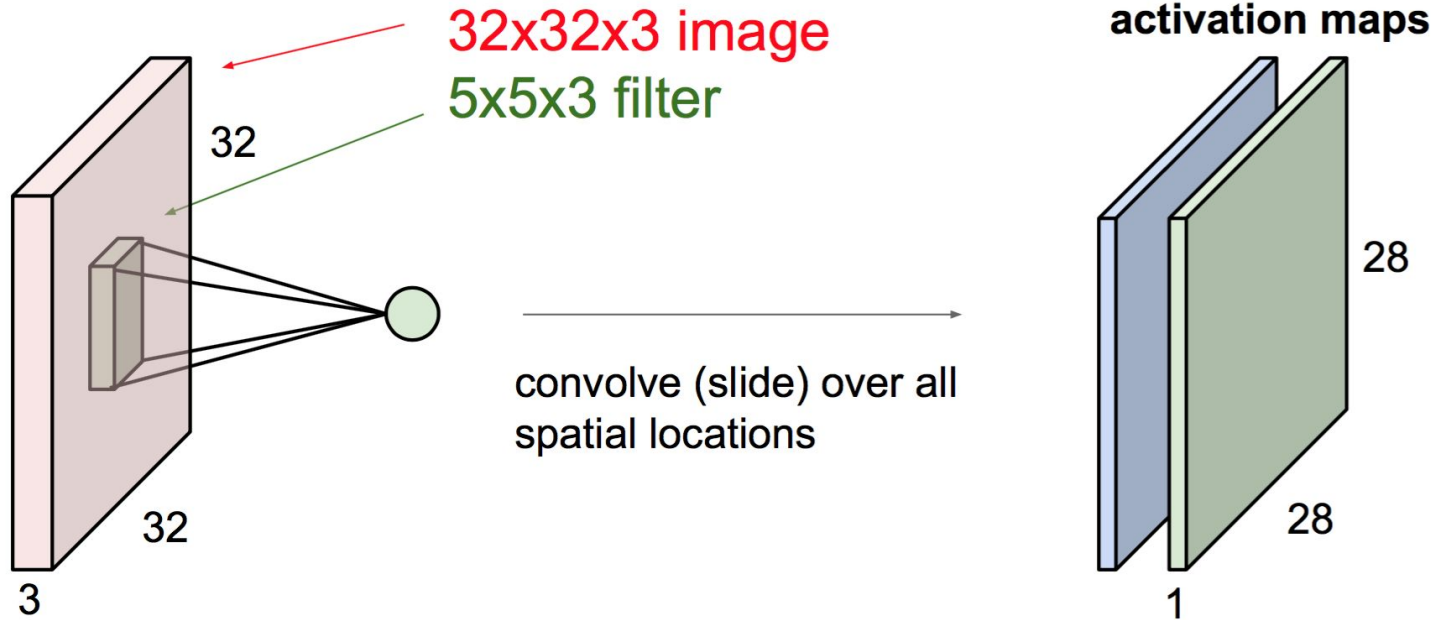
# Convolutional Neural Networks

## Convolutional Layer



# Convolutional Neural Networks


## Convolutional Layer





# Convolutional Neural Networks

## Convolutional Layer




We stack these up to get a “new image” of size 28x28x6!

# Convolutional Neural Networks

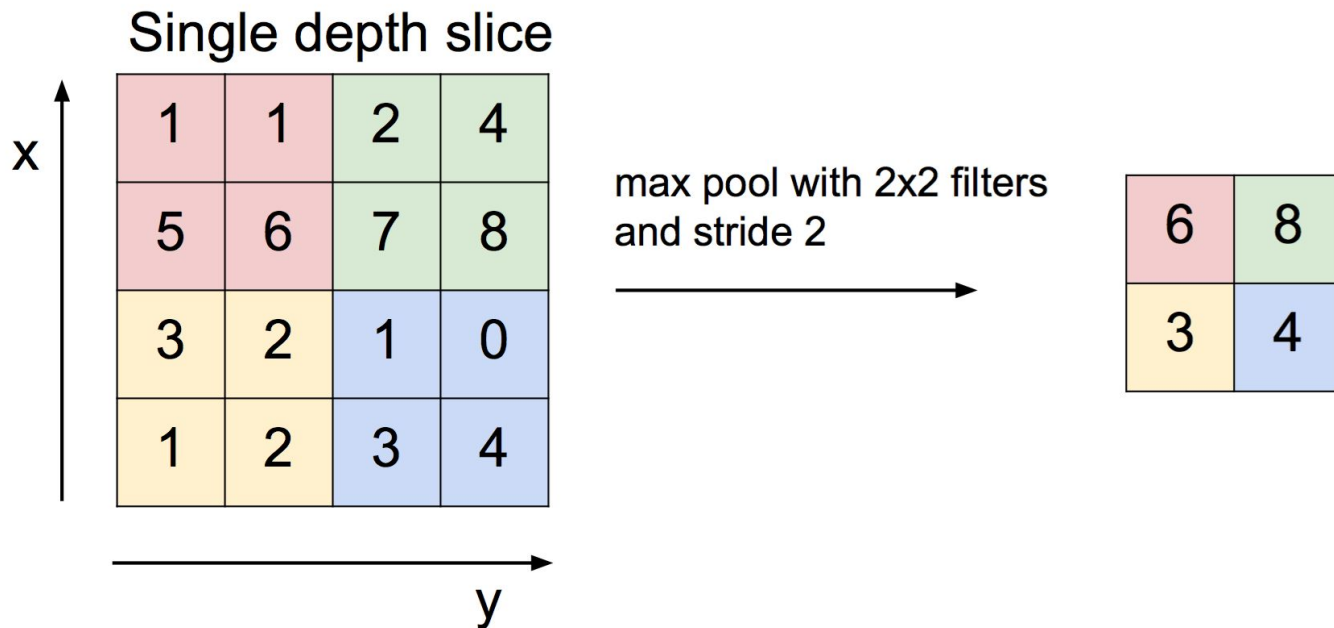
## Pooling Layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



# Convolutional Neural Networks

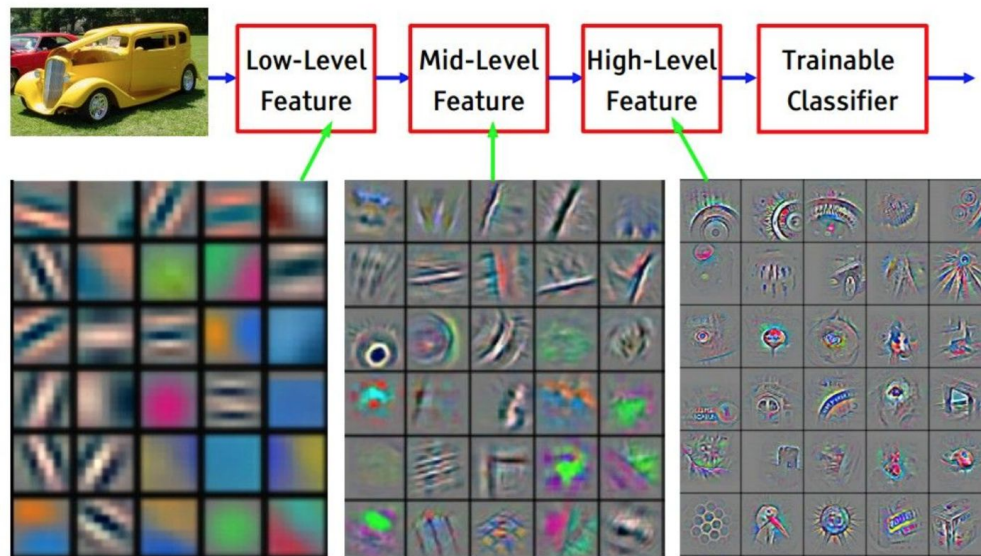
## Pooling Layer



# Convolutional Neural Networks

## Feature extraction

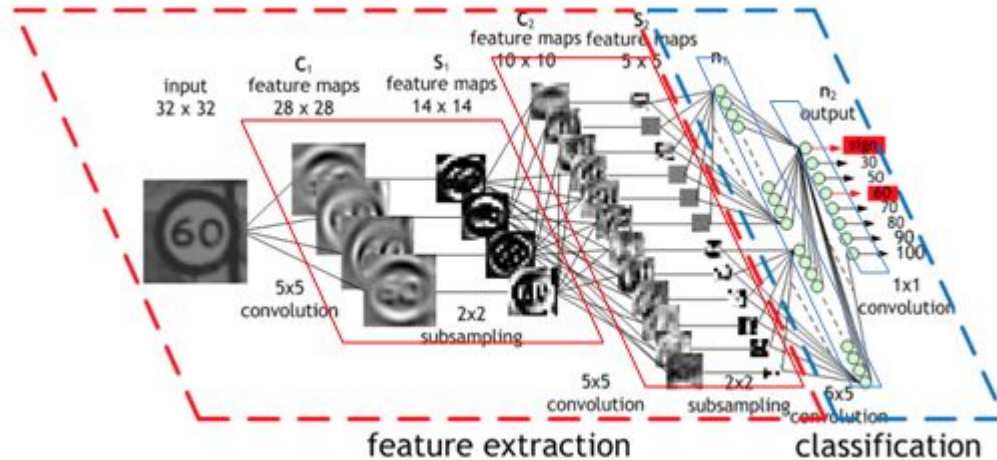
- › stack multiple Convolutional and Pooling Layer together



# Convolutional Neural Networks

## Fully Connected

- › normal neural network
- › classification through extracted features



# Convolutional Neural Networks

## Fully Connected

- › AlexNet

- › 224x224x3 images as input
- › 150.528 input values
- › feature extraction → 4096 important values

# Convolutional Neural Networks

## CNN in TensorFlow

*# convolutional layer*

```
conv = tf.layers.conv2d(input, 32, [3, 3], padding="same", activation=tf.nn.relu)
```

*# max-pooling layer*

```
max_pool = tf.layers.max_pooling2d(input, [2, 2], [2, 2])
```

*# fully connected*

```
fc = tf.layers.dense(input, 1024, activation=tf.nn.relu)
```

# Convolutional Neural Networks

CNN in TensorFlow

let's recognize some handwritten  
digits...

03\_cnn\_mnist.ipynb



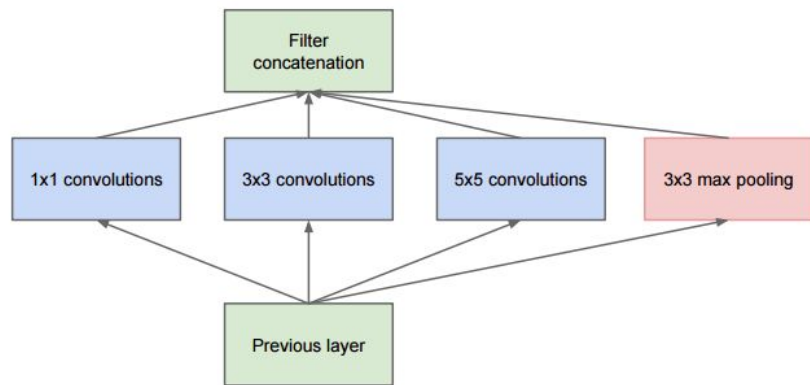
# Convolutional Neural Networks

## Inception

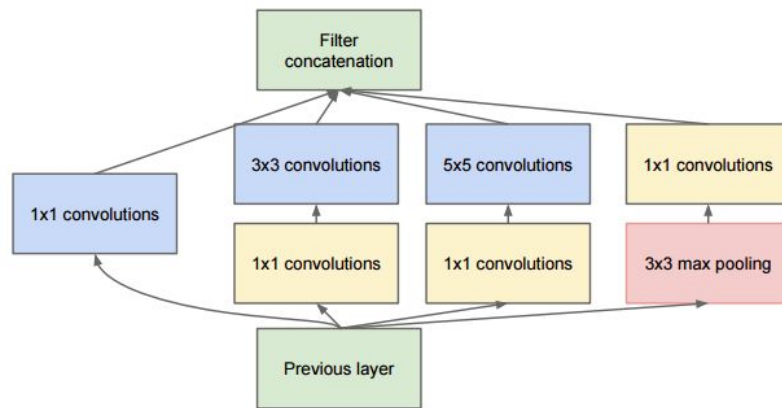
- › special architecture of very deep CNN
  - › very efficient training
  - › high performance
- › inception modules → parallel convolutions
- › ILSVRC 2012: Inception-ResNet-V2 → 95,3 % Accuracy

# Convolutional Neural Networks

## Inception



(a) Inception module, naïve version



(b) Inception module with dimension reductions


# Convolutional Neural Networks

## Inception-ResNet-v2

Inception Resnet V2 Network



Compressed View



# Convolutional Neural Networks

## Inception-ResNet-v2



# Convolutional Neural Networks

## CNN Model Zoo

- › state-of-the-art CNN
- › implemented in TF-slim → lightweight high-level API of TensorFlow
- › simple training and retraining
- › <https://github.com/tensorflow/models/tree/master/slim>